Project choises

The project contains two main folders:

- 'src' folder for all the source code.
- 'output' folder for *.txt* files that contains the **ping** command output for each ip address inserted in the program.

NetPerf

This project is written in python because this language provides useful modules such as matplotlib, numpy and pingparsing. These modules are used to manage charts, arrays and to parse the output of the ping command. *NetPerf* is compatible with both Linux and Windows operating system.

Instruction

Having GNU Make for building projects will simplify the execution of the program. Otherwise it is possible to run the program manually. This project requires some modules as previously stated. First of all pip module installer for python needs to be installed and added to the PATH. On linux if you have python already install type `sudo apt install python3-pip` .

For windows psping is required.

With Make

```
$ make install
$ make exe


        # or


$ make all
```

Without Make

```
$ python3 -m pip install pingparsing
$ python3 -m pip install matplotlib
$ python3 -m pip install numpy

$ mkdir output
$ python3 -B ./src/netperf.py
```

Once the program is running, it will ask to insert the hostname of which the **ping** command is going to run (notice that you can skip this input by simply pressing *Enter* and the default hostname will be

"lon.speedtest.clouvider.net") and if you are in linux it will ask also to decide whether the user wants to run the commands in *sudo* mode. In windows it is better practise to open the terminal in administrator mode in order to solve authorization problems. Then the program will run some ping commands in sequence and the traceroute command (tracert for windows) to find the correct amount of links crossed. Another ping command is prompted in order to save the output. At the end the program will loop other ping commands with increasing payload of packets each iteration up to 1472 bytes and steps of 16 bytes (number of packets sent is 20 each time). The last part of the program could take more than 10 minutes because the **ping** command is a time consuming process expecially if used for 90 iterations. In windows **psping** remove this time issue.

Example NetPerf

Considering:

- hostname = 'lon.speedtest.clouvider.net'

- Finding total hops -->

  - Linux:
    - command **ping** with params: *'-c 1'* to send only 1 packet, *'-t n'* to set max TTL to value *n* and *'-s 1'* to set packet size to 1 byte
    - command **traceroute** to verify the correct links crossed during the **ping**
  - Windows:
    - command **ping** with params: *'-n 1'* to send only 1 packet, *'-i n'* to set max TTL to value *n* and *'-l 1'* to set packet size to 1 byte
    - command **tracert** to verify the correct links crossed during the **ping**

- Default **ping** to save output to *.txt* -->

  - Linux:
    - command **ping** with params: *'-c 10'* to send 10 packets, *'-t 64'* to set max TTL to value *64* and *'-s 64'* to set packet size to 64 bytes
  - Windows:
    - command **ping** with params: *'-n 10'* to send 10 packets, *'-i 64'* to set max TTL to value *64* and *'-l 64'* to set packet size to 64 bytes

- **Ping** loop to get the network stats -->

  - Set a list of steps of bytes starting from 64 up to 1472 with steps of 16 bytes.
  - Linux:
    - command **ping** with params: *'-c 20'* to send 20 packets, *'-t 64'* to set max TTL to default value *64* and *'-s steps[i]'* to set packet size to *steps[i]* bytes
  - Windows:
    - command **ping** with params: *'-n 20'* to send 20 packets, *'-i 0'* to set waiting time to *0* and *'-l steps[i]'* to set packet size to *steps[i]* bytes

LINUX

```
$ make exe

$ if [ ! -d "./output" ]; then mkdir output; fi
$ python3 -B ./src/netperf.py
$ Insert hostname, skip by pressing enter. The default hostname is
$ 'lon.speedtest.clouvider.net'
$ --> lon.speedtest.clouvider.net
$ IP: 5.180.211.133
$ Sudo (Y/N) Y                    // this input is present only in linux

$------------Finding amount of hops to the host------------
$ Running command: sudo ping -c 1 -t 30 -s 1 5.180.211.133
$ [sudo] password for dave:
$ Running command: sudo ping -c 1 -t 29 -s 1 5.180.211.133
$ Running command: sudo ping -c 1 -t 28 -s 1 5.180.211.133
$ Running command: sudo ping -c 1 -t 27 -s 1 5.180.211.133
$ Running command: sudo ping -c 1 -t 26 -s 1 5.180.211.133
$ Running command: sudo ping -c 1 -t 25 -s 1 5.180.211.133
$ Running command: sudo ping -c 1 -t 24 -s 1 5.180.211.133
$ Running command: sudo ping -c 1 -t 23 -s 1 5.180.211.133
$ Running command: sudo ping -c 1 -t 22 -s 1 5.180.211.133
$ Running command: sudo ping -c 1 -t 21 -s 1 5.180.211.133
$ Running command: sudo ping -c 1 -t 20 -s 1 5.180.211.133
$ Running command: sudo ping -c 1 -t 19 -s 1 5.180.211.133
$ ------------------------------------------------------------
$ Running command: sudo ping -c 10 -t 64 -s 64 5.180.211.133
$ PING 5.180.211.133 (5.180.211.133) 64(92) bytes of data.
$ 72 bytes from 5.180.211.133: icmp_seq=1 ttl=54 time=89.9 ms
$ 72 bytes from 5.180.211.133: icmp_seq=2 ttl=54 time=112 ms
$ 72 bytes from 5.180.211.133: icmp_seq=3 ttl=54 time=36.7 ms
$ 72 bytes from 5.180.211.133: icmp_seq=4 ttl=54 time=53.2 ms
$ 72 bytes from 5.180.211.133: icmp_seq=5 ttl=54 time=73.9 ms
$ 72 bytes from 5.180.211.133: icmp_seq=6 ttl=54 time=55.4 ms
$ 72 bytes from 5.180.211.133: icmp_seq=7 ttl=54 time=118 ms
$ 72 bytes from 5.180.211.133: icmp_seq=8 ttl=54 time=100 ms
$ 72 bytes from 5.180.211.133: icmp_seq=9 ttl=54 time=60.4 ms
$ 72 bytes from 5.180.211.133: icmp_seq=10 ttl=54 time=81.7 ms
$
$ --- 5.180.211.133 ping statistics ---
$ 10 packets transmitted, 10 received, 0% packet loss, time 9019ms
$ rtt min/avg/max/mdev = 36.665/78.131/118.235/25.639 ms
$ ------------------------------------
$ Running command: traceroute 5.180.211.133
$ traceroute to 5.180.211.133 (5.180.211.133), 30 hops max, 60 byte packets
$  1  myfastgate.lan (192.168.1.254)  4.492 ms  4.456 ms  4.445 ms
$  2  10.5.3.124 (10.5.3.124)  28.441 ms  16.458 ms  28.402 ms
$  3  10.5.105.106 (10.5.105.106)  12.125 ms  12.372 ms  16.415 ms
$  4  10.250.137.205 (10.250.137.205)  16.405 ms 10.250.137.201 (10.250.137.201)  16
$  5  10.254.20.41 (10.254.20.41)  23.915 ms  23.905 ms 10.254.20.45 (10.254.20.45)
$  6  93-63-100-105.ip27.fastwebnet.it (93.63.100.105)  23.884 ms 93-63-100-113.ip27
```

```
$  7  93-63-100-141.ip27.fastwebnet.it (93.63.100.141)  15.298 ms  20.303 ms 93-63-1
$  8  i63lgv-005-ae5.bb.ip-plus.net (193.5.122.134)  16.955 ms  16.939 ms  20.570 ms
$  9  i71lzw-015-ae4.bb.ip-plus.net (138.187.129.162)  20.561 ms  20.550 ms  17.244
$ 10  i79zhh-005-ae21.bb.ip-plus.net (138.187.129.75)  17.214 ms  17.203 ms  40.430
$ 11  i79tix-025-ae10.bb.ip-plus.net (138.187.129.174)  40.399 ms  40.382 ms  40.342
$ 12  100ge0-35.core2.zrh3.he.net (72.52.92.130)  40.323 ms  102.601 ms  102.558 ms
$ 13  100ge0-35.core2.gva1.he.net (184.104.193.134)  102.537 ms  102.512 ms *
$ 14  * * *
$ 15  linx-lon1.thn2.peering.clouvider.net (195.66.227.14)  101.388 ms  101.343 ms
$ 16  * * *
$ 17  h185-42-222-17.reverse.clouvider.net (185.42.222.17)  102.647 ms  102.628 ms *
$ 18  185.245.80.45 (185.245.80.45)  102.521 ms  102.503 ms  102.485 ms
$ 19  185.245.80.1 (185.245.80.1)  98.261 ms  63.026 ms  62.979 ms
$ 20  5.180.211.133 (5.180.211.133)  62.960 ms  62.941 ms 194.35.234.205 (194.35.234
$ ------------------------------------
$ Route is correct
$ Hops to host: 20
$ Number of links crossed: 40
$ +++++++++++++++++++++++++++++++++++++++++++++
$ Running netperf with packets of size 512 bits
$ Running command: sudo ping -c 20 -t 64 -s 64 5.180.211.133

        // ... iteration of the ping command to get the network stats
```
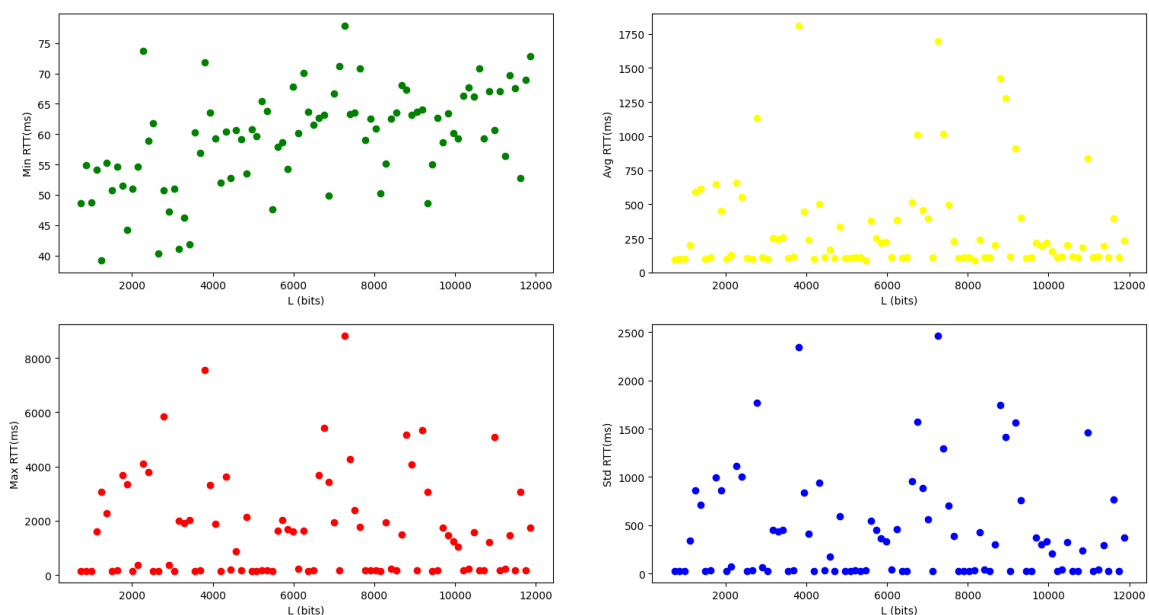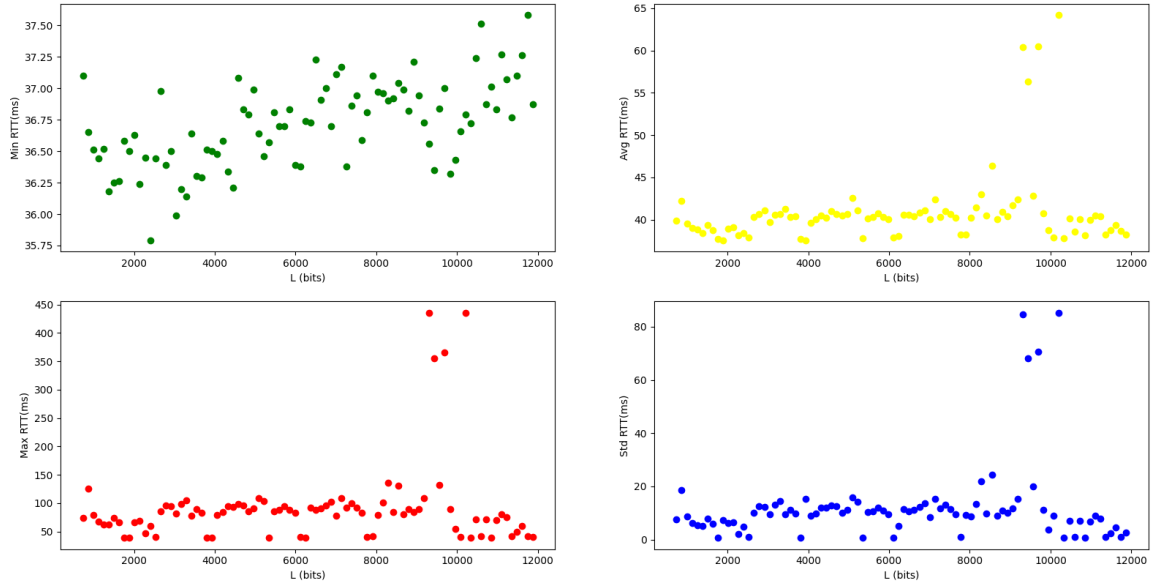
Results

In the previous example the number of links crossed is 32, and it is verified by the *traceroute* or *tracert* command. The RTT data collected from the last part of the program is displayable in these graphs:

Linux:

Windows:

From min-RTT graph we can use the **numpy.polyfit** method to get the slope of the function. We can call this coefficient *'a'* (notice that *'a'* has to be converted from *ms/bit* to *s/bit*). This slope is used to estimate the value of the *throughput* and the *throughput bottleneck* with the following formulas:

$$S = \frac{n_{links}}{a}$$

$$S_{bottleneck} = \frac{2}{a}$$

In this specific example the value of the throughputs are:

- Linux:

$$S = 590440329.6 \; bits/s$$

$$S_b = 33222016.48 \; bits/s$$

- Windows:

$$S = 578657266.07 \; bits/s$$

$$S_b = 32147625.89 \; bits/s$$

With **numpy.polyfit** we can also calculate the time of propagation *T* of the signal through the links.

- Linux:

$$T = 38.55 \; ms$$

- Windows:

$$T = 36.32 \; ms$$

We can see that the throughput is in the order of magnitude of tens of Mbit/s in optimal conditions, meaning that all the links in the connection have the same throughput and the packets queue time in each link is approximately *0s*. If there is a single link which has the lowest throughput among all of the links crossed, it works as a bottleneck limiting the thoughput (notice that this link as to be considered two times because the packets need to return through the same route).

The time of propagation does not depend from the network speed but from the actual speed of the propagation of the signal, so it is governed by physics laws that prevents the signal from being faster than speed of light. If we ping to an host located in America the minimum time of propagation is around *120 ms*, otherwise if we ping to an host located in Europe the *T* can be close to *30 ms*. Notice that this time depends on the type of connection, wired connection performs better.